

Problem Statement: Houses are one of the necessary need of each and every person around the globe and therefore housing and real estate market is one of the markets which is one of the major contributors in the world's economy. It is a very large market and there are various companies working in the domain. Data science comes as a very important tool to solve problems in the domain to help the companies increase their overall revenue, profits, improving their marketing strategies and focusing on changing trends in house sales and purchases. Predictive modelling, Market mix modelling, recommendation systems are some of the machine learning techniques used for achieving the business goals for housing companies. Our problem is related to one such housing company. A US-based housing company named Surprise Housing has decided to enter the Australian market. The company uses data analytics to purchase houses at a price below their actual values and flip them at a higher price. For the same purpose, the company has collected a data set from the sale of houses in Australia. The data is provided in the CSV file below. The company is looking at prospective properties to buy houses to enter the market. You are required to build a model using Machine Learning in order to predict the actual value of the prospective properties and decide whether to invest in them or not. For this company wants to know: • Which variables are important to predict the price of variable? • How do these variables describe the price of the house

import libraries

```
In [1]: 1 import pandas as pd
        2 import numpy as np
        3 import seaborn as sns
        4 import matplotlib.pyplot as plt
        5 from scipy.stats import norm, skew
        6 import warnings
        7 warnings.filterwarnings('ignore')
```

```
In [2]: 1 df=pd.read_csv(r'C:\Users\polasasuresh\Downloads\Project-Housing--2---1- (2)
2 df
```

Out[2]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	l
0	127	120	RL	NaN	4928	Pave	NaN	IR1		Lvl
1	889	20	RL	95.0	15865	Pave	NaN	IR1		Lvl
2	793	60	RL	92.0	9920	Pave	NaN	IR1		Lvl
3	110	20	RL	105.0	11751	Pave	NaN	IR1		Lvl
4	422	20	RL	NaN	16635	Pave	NaN	IR1		Lvl
...	...	...	...	...	...	...	...	...		...
1163	289	20	RL	NaN	9819	Pave	NaN	IR1		Lvl
1164	554	20	RL	67.0	8777	Pave	NaN	Reg		Lvl
1165	196	160	RL	24.0	2280	Pave	NaN	Reg		Lvl
1166	31	70	C (all)	50.0	8500	Pave	Pave	Reg		Lvl
1167	617	60	RL	NaN	7861	Pave	NaN	IR1		Lvl

1168 rows × 81 columns

```
In [3]: 1 df_Test=pd.read_csv(r'C:\Users\polasasuresh\Downloads\Project-Housing--2---1
2 df_Test
```

Out[3]:

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	LandContour	l
0	337	20	RL	86.0	14157	Pave	NaN	IR1		HLS
1	1018	120	RL	NaN	5814	Pave	NaN	IR1		Lvl
2	929	20	RL	NaN	11838	Pave	NaN	Reg		Lvl
3	1148	70	RL	75.0	12000	Pave	NaN	Reg		Bnk
4	1227	60	RL	86.0	14598	Pave	NaN	IR1		Lvl
...	...	...	...	...	...	...	...	...		...
287	83	20	RL	78.0	10206	Pave	NaN	Reg		Lvl
288	1048	20	RL	57.0	9245	Pave	NaN	IR2		Lvl
289	17	20	RL	NaN	11241	Pave	NaN	IR1		Lvl
290	523	50	RM	50.0	5000	Pave	NaN	Reg		Lvl
291	1379	160	RM	21.0	1953	Pave	NaN	Reg		Lvl

292 rows × 80 columns

```
In [4]: 1 df.shape
```

Out[4]: (1168, 81)

In [5]: 1 df\_Test.shape

Out[5]: (292, 80)

There are 1168 rows and 81 columns

1 From looking at the both sets, we can see that the only difference in features is "Sale Price". This makes sense because we are trying to predict it!

In [6]: 1 df.dtypes

Out[6]: Id int64  
 MSSubClass int64  
 MSZoning object  
 LotFrontage float64  
 LotArea int64  
 ...  
 MoSold int64  
 YrSold int64  
 SaleType object  
 SaleCondition object  
 SalePrice int64  
 Length: 81, dtype: object

In [7]: 1 df.columns

Out[7]: Index(['Id', 'MSSubClass', 'MSZoning', 'LotFrontage', 'LotArea', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities', 'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2', 'BldgType', 'HouseStyle', 'OverallQual', 'OverallCond', 'YearBuilt', 'YearRemodAdd', 'RoofStyle', 'RoofMatl', 'Exterior1st', 'Exterior2nd', 'MasVnrType', 'MasVnrArea', 'ExterQual', 'ExterCond', 'Foundation', 'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinSF1', 'BsmtFinType2', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'Heating', 'HeatingQC', 'CentralAir', 'Electrical', '1stFlrSF', '2ndFlrSF', 'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath', 'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'KitchenQual', 'TotRmsAbvGrd', 'Functional', 'Fireplaces', 'FireplaceQu', 'GarageType', 'GarageYrBlt', 'GarageFinish', 'GarageCars', 'GarageArea', 'GarageQual', 'GarageCond', 'PavedDrive', 'WoodDeckSF', 'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea', 'PoolQC', 'Fence', 'MiscFeature', 'MiscVal', 'MoSold', 'YrSold', 'SaleType', 'SaleCondition', 'SalePrice'], dtype='object')

There are 81 diifferent columns present in the dataset

In [8]:

```

1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1168 entries, 0 to 1167
Data columns (total 81 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Id                    1168 non-null   int64
 1   MSSubClass            1168 non-null   int64
 2   MSZoning              1168 non-null   object
 3   LotFrontage          954 non-null    float64
 4   LotArea              1168 non-null   int64
 5   Street               1168 non-null   object
 6   Alley               77 non-null     object
 7   LotShape             1168 non-null   object
 8   LandContour          1168 non-null   object
 9   Utilities            1168 non-null   object
10  LotConfig            1168 non-null   object
11  LandSlope            1168 non-null   object
12  Neighborhood          1168 non-null   object
13  Condition1           1168 non-null   object
14  Condition2           1168 non-null   object
15  BldgType             1168 non-null   object
16  HouseStyle           1168 non-null   object
17  OverallQual          1168 non-null   int64
18  OverallCond          1168 non-null   int64
19  YearBuilt            1168 non-null   int64
20  YearRemodAdd         1168 non-null   int64
21  RoofStyle            1168 non-null   object
22  RoofMatl            1168 non-null   object
23  Exterior1st         1168 non-null   object
24  Exterior2nd         1168 non-null   object
25  MasVnrType          1161 non-null   object
26  MasVnrArea          1161 non-null   float64
27  ExterQual           1168 non-null   object
28  ExterCond           1168 non-null   object
29  Foundation          1168 non-null   object
30  BsmtQual            1138 non-null   object
31  BsmtCond            1138 non-null   object
32  BsmtExposure        1137 non-null   object
33  BsmtFinType1        1138 non-null   object
34  BsmtFinSF1          1168 non-null   int64
35  BsmtFinType2        1137 non-null   object
36  BsmtFinSF2          1168 non-null   int64
37  BsmtUnfSF           1168 non-null   int64
38  TotalBsmtSF         1168 non-null   int64
39  Heating             1168 non-null   object
40  HeatingQC           1168 non-null   object
41  CentralAir          1168 non-null   object
42  Electrical           1168 non-null   object
43  1stFlrSF            1168 non-null   int64
44  2ndFlrSF            1168 non-null   int64
45  LowQualFinSF        1168 non-null   int64
46  GrLivArea           1168 non-null   int64
47  BsmtFullBath        1168 non-null   int64
48  BsmtHalfBath        1168 non-null   int64
49  FullBath            1168 non-null   int64

```

```

50 HalfBath          1168 non-null    int64
51 BedroomAbvGr     1168 non-null    int64
52 KitchenAbvGr     1168 non-null    int64
53 KitchenQual       1168 non-null    object
54 TotRmsAbvGrd     1168 non-null    int64
55 Functional        1168 non-null    object
56 Fireplaces        1168 non-null    int64
57 FireplaceQu       617 non-null     object
58 GarageType        1104 non-null    object
59 GarageYrBlt       1104 non-null    float64
60 GarageFinish      1104 non-null    object
61 GarageCars        1168 non-null    int64
62 GarageArea        1168 non-null    int64
63 GarageQual        1104 non-null    object
64 GarageCond        1104 non-null    object
65 PavedDrive        1168 non-null    object
66 WoodDeckSF        1168 non-null    int64
67 OpenPorchSF       1168 non-null    int64
68 EnclosedPorch     1168 non-null    int64
69 3SsnPorch         1168 non-null    int64
70 ScreenPorch       1168 non-null    int64
71 PoolArea          1168 non-null    int64
72 PoolQC            7 non-null       object
73 Fence             237 non-null     object
74 MiscFeature       44 non-null      object
75 MiscVal           1168 non-null    int64
76 MoSold            1168 non-null    int64
77 YrSold            1168 non-null    int64
78 SaleType          1168 non-null    object
79 SaleCondition     1168 non-null    object
80 SalePrice         1168 non-null    int64
dtypes: float64(3), int64(35), object(43)
memory usage: 543.0+ KB

```

There are 35 integer type of data and 43 object type of data and 3 float type of data present in the dataset

## EDA

In [9]: 1 df.isnull().sum()

```


Out[9]: Id          0
MSSubClass         0
MSZoning           0
LotFrontage       214
LotArea           0
...
MoSold            0
YrSold            0
SaleType          0
SaleCondition     0
SalePrice         0
Length: 81, dtype: int64

```

There are some null values present in the data

```
In [10]: 1 sns.heatmap(df.isnull())
```

```
Out[10]: <AxesSubplot:>
```



## 2. Analyzing the Test Variable (Sale Price)

```
1
2 Let's check out the most interesting feature in this study: Sale Price.
   Important Note: This data is from Ames, Iowa. The location is extremely
   correlated with Sale Price. (I had to take a double-take at a point, since
   I consider myself a house-browsing enthusiast)
```

```
In [12]: 1 # Getting Description
          2 df['SalePrice'].describe()
```

```
Out[12]: count      1168.000000
          mean      181477.005993
          std       79105.586863
          min       34900.000000
          25%      130375.000000
          50%      163995.000000
          75%      215000.000000
          max       755000.000000
          Name: SalePrice, dtype: float64
```

With an average house price of \$180921, it seems like I should relocated to Iowa!

```
In [ ]: 1
```

```
In [ ]: 1
```

In [ ]:

1

In [ ]:

1

In [ ]:

1

In [ ]:

1

### 3. Multivariable Analysis

Let's check out all the variables! There are two types of features in housing data, categorical and numerical.

Categorical data is just like it sounds. It is in categories. It isn't necessarily linear, but it follows some kind of pattern. For example, take a feature of "Downtown". The response is either "Near", "Far", "Yes", and "No". Back then, living in downtown usually meant that you couldn't afford to live in uptown. Thus, it could be implied that downtown establishments cost less to live in. However, today, that is not the case. (Thank you, hipsters!) So we can't really establish any particular order of response to be "better" or "worse" than the other.

Numerical data is data in number form. (Who could have thought!) These features are in a linear relationship with each other. For example, a 2,000 square foot place is 2 times "bigger" than a 1,000 square foot place. Plain and simple. Simple and clean.

In [14]:

```
1 # Checking Categorical Data
2 df.select_dtypes(include=['object']).columns
```

```
Out[14]: Index(['MSZoning', 'Street', 'Alley', 'LotShape', 'LandContour', 'Utilities',
               'LotConfig', 'LandSlope', 'Neighborhood', 'Condition1', 'Condition2',
               'BldgType', 'HouseStyle', 'RoofStyle', 'RoofMat1', 'Exterior1st',
               'Exterior2nd', 'MasVnrType', 'ExterQual', 'ExterCond', 'Foundation',
               'BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFinType2',
               'Heating', 'HeatingQC', 'CentralAir', 'Electrical', 'KitchenQual',
               'Functional', 'FireplaceQu', 'GarageType', 'GarageFinish', 'GarageQual',
               'GarageCond', 'PavedDrive', 'PoolQC', 'Fence', 'MiscFeature',
               'SaleType', 'SaleCondition'],
              dtype='object')
```

In [15]:

```
1 # Checking Numerical Data
2 df.select_dtypes(include=['int64', 'float64']).columns
```

```
Out[15]: Index(['Id', 'MSSubClass', 'LotFrontage', 'LotArea', 'OverallQual',
               'OverallCond', 'YearBuilt', 'YearRemodAdd', 'MasVnrArea', 'BsmtFinSF1',
               'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', '1stFlrSF', '2ndFlrSF',
               'LowQualFinSF', 'GrLivArea', 'BsmtFullBath', 'BsmtHalfBath', 'FullBath',
               'HalfBath', 'BedroomAbvGr', 'KitchenAbvGr', 'TotRmsAbvGrd',
               'Fireplaces', 'GarageYrBlt', 'GarageCars', 'GarageArea', 'WoodDeckSF',
               'OpenPorchSF', 'EnclosedPorch', '3SsnPorch', 'ScreenPorch', 'PoolArea',
               'MiscVal', 'MoSold', 'YrSold', 'SalePrice'],
              dtype='object')
```

```
In [16]: 1 cat = len(df.select_dtypes(include=['object']).columns)
          2 num = len(df.select_dtypes(include=['int64', 'float64']).columns)
          3 print('Total Features: ', cat, 'categorical', '+',
          4       num, 'numerical', '=', cat+num, 'features')
```

Total Features: 43 categorical + 38 numerical = 81 features

With 81 features, how could we possibly tell which feature is most related to house prices? Good thing we have a correlation matrix. Let's do it!



In [17]:


1 df.corr()

Out[17]:

	Id	MSSubClass	LotFrontage	LotArea	OverallQual	OverallCond	YearBuilt
Id	1.000000	0.004259	-0.006629	-0.029212	-0.036965	0.039761	-0.016942
MSSubClass	0.004259	1.000000	-0.365220	-0.124151	0.070462	-0.056978	0.023988
LotFrontage	-0.006629	-0.365220	1.000000	0.557257	0.247809	-0.053345	0.118554
LotArea	-0.029212	-0.124151	0.557257	1.000000	0.107188	0.017513	0.005506
OverallQual	-0.036965	0.070462	0.247809	0.107188	1.000000	-0.083167	0.575800
OverallCond	0.039761	-0.056978	-0.053345	0.017513	-0.083167	1.000000	-0.377731
YearBuilt	-0.016942	0.023988	0.118554	0.005506	0.575800	-0.377731	1.000000
YearRemodAdd	-0.018590	0.056618	0.096050	0.027228	0.555945	0.080669	0.596050
MasVnrArea	-0.060652	0.027868	0.202225	0.121448	0.409163	-0.137882	0.322225
BsmtFinSF1	0.003868	-0.052236	0.247780	0.221851	0.219643	-0.028810	0.227927
BsmtFinSF2	0.005269	-0.062403	0.002514	0.056656	-0.040893	0.044336	-0.025144
BsmtUnfSF	-0.019494	-0.134170	0.123943	0.006600	0.308676	-0.146384	0.151949
TotalBsmtSF	-0.013812	-0.214042	0.386261	0.259733	0.528285	-0.162481	0.386261
1stFlrSF	0.009647	-0.227927	0.448186	0.312843	0.458758	-0.134420	0.279279
2ndFlrSF	-0.029671	0.300366	0.099250	0.059803	0.316624	0.036668	0.010993
LowQualFinSF	-0.070180	0.053737	0.007885	-0.001915	-0.039295	0.041877	-0.189180
GrLivArea	-0.024325	0.086448	0.410414	0.281360	0.599700	-0.065006	0.190414
BsmtFullBath	0.023027	0.004556	0.104255	0.142387	0.101732	-0.039680	0.164255
BsmtHalfBath	-0.043572	0.008207	0.001528	0.059282	-0.030702	0.091016	-0.028207
FullBath	-0.015187	0.140807	0.189321	0.123197	0.548824	-0.171931	0.479321
HalfBath	-0.028512	0.168423	0.053168	0.007271	0.296134	-0.052125	0.248423
BedroomAbvGr	0.009376	-0.013283	0.264010	0.117351	0.099639	0.028393	-0.084010
KitchenAbvGr	0.001216	0.283506	-0.002890	-0.013075	-0.178220	-0.076047	-0.162890
TotRmsAbvGrd	-0.001613	0.051179	0.351969	0.184546	0.432579	-0.039952	0.091969
Fireplaces	-0.024175	-0.035792	0.262076	0.285983	0.390067	-0.013632	0.132076
GarageYrBlt	-0.000469	0.077630	0.061101	-0.034981	0.541719	-0.318278	0.821101
GarageCars	0.007549	-0.027639	0.276798	0.158313	0.596322	-0.161996	0.527639
GarageArea	0.010048	-0.092408	0.344908	0.195162	0.566782	-0.126021	0.474908
WoodDeckSF	-0.027498	-0.022609	0.101751	0.216720	0.227137	0.012290	0.201751
OpenPorchSF	-0.013642	0.017468	0.167092	0.093080	0.341030	-0.024899	0.193080
EnclosedPorch	0.004885	-0.004252	0.023118	-0.007446	-0.098374	0.056074	-0.373118
3SsnPorch	-0.021773	-0.043210	0.059508	0.025794	0.045919	0.040476	0.039508
ScreenPorch	0.005169	-0.013291	0.033111	0.025256	0.059387	0.069463	-0.051329
PoolArea	0.065832	0.009583	0.223429	0.097107	0.072247	-0.003603	0.009583


		<b>Id</b>	<b>MSSubClass</b>	<b>LotFrontage</b>	<b>LotArea</b>	<b>OverallQual</b>	<b>OverallCond</b>	<b>Year</b>
	<b>MiscVal</b>	0.001304	-0.023503	-0.004559	0.051679	-0.025786	0.075178	-0.030
	<b>MoSold</b>	0.023479	-0.016015	0.025046	0.015141	0.090638	0.005519	0.03
	<b>YrSold</b>	-0.008853	-0.038595	-0.004296	-0.035399	-0.048759	0.055517	-0.01
	<b>SalePrice</b>	-0.023897	-0.060775	0.341294	0.249499	0.789185	-0.065642	0.51

38 rows × 38 columns



```
In [18]: 1 plt.figure(figsize=(20,10))
         2 sns.heatmap(df.corr(),annot=True)
```


Out[18]: <AxesSubplot:>



```

In [20]: 1 # Top 10 Heatmap
          2 k = 10 #number of variables for heatmap
          3 cols = df.corr().nlargest(k, 'SalePrice')['SalePrice'].index
          4 cm = np.corrcoef(df[cols].values.T)
          5 sns.set(font_scale=1.25)
          6 hm = sns.heatmap(cm, cbar=True, annot=True, square=True, fmt='.2f', annot_kw
          7 plt.show()

```



```
In [21]: 1 most_corr = pd.DataFrame(cols)
          2 most_corr.columns = ['Most Correlated Features']
          3 most_corr
```


Out[21]:

Most Correlated Features	
0	SalePrice
1	OverallQual
2	GrLivArea
3	GarageCars
4	GarageArea
5	TotalBsmtSF
6	1stFlrSF
7	FullBath
8	TotRmsAbvGrd
9	YearBuilt

Well, the most correlated feature to Sale Price is... Sale Price?!? Of course. For the other 9, they are as listed. Here is a short description of each. (Thank you, data\_description.txt!)

OverallQual: Rates the overall material and finish of the house (1 = Very Poor, 10 = Very Excellent)  
GrLivArea: Above grade (ground) living area square feet  
GarageCars: Size of garage in car capacity  
GarageArea: Size of garage in square feet  
TotalBsmtSF: Total square feet of basement area  
1stFlrSF: First Floor square feet  
FullBath: Full bathrooms above grade  
TotRmsAbvGrd: Total rooms above grade (does not include bathrooms)  
YearBuilt: Original construction date  
Let's take a look at how each relates to Sale Price and do some pre-cleaning on each feature if necessary.


```
In [22]: 1 # Overall Quality vs Sale Price
2 var = 'OverallQual'
3 data = pd.concat([df['SalePrice'], df[var]], axis=1)
4 f, ax = plt.subplots(figsize=(8, 6))
5 fig = sns.boxplot(x=var, y="SalePrice", data=data)
6 fig.axis(ymin=0, ymax=800000);
```



People pay more for better quality? Nothing new here. Let's move on.

```
In [23]: 1 # Living Area vs Sale Price
         2 sns.jointplot(x=df['GrLivArea'], y=df['SalePrice'], kind='reg')
```

Out[23]: <seaborn.axisgrid.JointGrid at 0xfa045b0>




It makes sense that people would pay for the more living area. What doesn't make sense is the two datapoints in the bottom-right of the plot.

We need to take care of this! What we will do is remove these outliers manually.

```
In [24]: 1 # Removing outliers manually (Two points in the bottom right)
         2 df = df.drop(df[(df['GrLivArea']>4000)
         3                 & (df['SalePrice']<300000)].index).reset_index(drop
```

```
In [25]: 1 # Living Area vs Sale Price  
        2 sns.jointplot(x=df['GrLivArea'], y=df['SalePrice'], kind='reg')
```


Out[25]: <seaborn.axisgrid.JointGrid at 0x1079d370>



Nice! We got a 0.02 point increase in the Pearson-R Score.

```
In [26]: 1 # Garage Area vs Sale Price
         2 sns.boxplot(x=df['GarageCars'], y=df['SalePrice'])
```

Out[26]: <AxesSubplot:xlabel='GarageCars', ylabel='SalePrice'>




4-car garages result in less Sale Price? That doesn't make much sense. Let's remove those outliers.

```
In [27]: 1 # Removing outliers manually (More than 4-cars, Less than $300k)
         2 df = df.drop(df[(df['GarageCars']>3)
         3               & (df['SalePrice']<300000)].index).reset_index(drop
```



```
In [28]: 1 # Garage Area vs Sale Price  
        2 sns.boxplot(x=df['GarageCars'], y=df['SalePrice'])
```


```
Out[28]: <AxesSubplot:xlabel='GarageCars', ylabel='SalePrice'>
```



That looks much better. Note: removal of data is totally discretionary and may or may not help in modeling. Use at your own preference.

```
In [29]: 1 # Garage Area vs Sale Price
         2 sns.jointplot(x=df['GarageArea'], y=df['SalePrice'], kind='reg')
```

Out[29]: <seaborn.axisgrid.JointGrid at 0x10d99838>




Again with the bottom two data-points. Let's remove those outliers.

```
In [30]: 1 # Removing outliers manually (More than 1000 sqft, Less than $300k)
         2 df= df.drop(df[(df['GarageArea']>1000)
         3                 & (df['SalePrice']<300000)].index).reset_index(drop
```

```
In [31]: 1 # Garage Area vs Sale Price  
        2 sns.jointplot(x=df['GarageArea'], y=df['SalePrice'], kind='reg')
```


Out[31]: <seaborn.axisgrid.JointGrid at 0x10d4e508>



Only 0.01 point Pearson-R Score increase, but looks much better!

```
In [33]: 1 # Basement Area vs Sale Price  
2 sns.jointplot(x=df['TotalBsmtSF'], y=df['SalePrice'], kind='reg')
```


Out[33]: <seaborn.axisgrid.JointGrid at 0x110ab328>



Everything looks fine here.

```
In [34]: 1 # First Floor Area vs Sale Price  
        2 sns.jointplot(x=df['1stFlrSF'], y=df['SalePrice'], kind='reg')
```


Out[34]: <seaborn.axisgrid.JointGrid at 0x10a79f28>



Looks good.


```
In [36]: 1 # Total Rooms vs Sale Price
          2 sns.boxplot(x=df['TotRmsAbvGrd'], y=df['SalePrice'])
```

Out[36]: <AxesSubplot:xlabel='TotRmsAbvGrd', ylabel='SalePrice'>



It seems like houses with more than 11 rooms come with a \$100k off coupon. It looks like an outlier but I'll let it slide.

```
In [37]: 1 # Total Rooms vs Sale Price
          2 var = 'YearBuilt'
          3 data = pd.concat([df['SalePrice'], df[var]], axis=1)
          4 f, ax = plt.subplots(figsize=(16, 8))
          5 fig = sns.boxplot(x=var, y="SalePrice", data=data)
          6 fig.axis(ymin=0, ymax=800000);
          7 plt.xticks(rotation=90);
```



Although it seems like house prices decrease with age, we can't be entirely sure. Is it because of inflation or stock market crashes? Let's leave the years alone.

## . Impute Missing Data and Clean Data

Important questions when thinking about missing data:

How prevalent is the missing data? Is missing data random or does it have a pattern? The answer to these questions is important for practical reasons because missing data can imply a reduction of the sample size. This can prevent us from proceeding with the analysis. Moreover, from a substantive perspective, we need to ensure that the missing data process is not biased and hiding an inconvenient truth.

Let's combine both training and test data into one dataset to impute missing values and do some cleaning.

In [40]:

```
1 # Combining Datasets
2 ntrain = df.shape[0]
3 ntest = df_Test.shape[0]
4 y_train = df.SalePrice.values
5 all_data = pd.concat((df, df_Test)).reset_index(drop=True)
6 all_data.drop(['SalePrice'], axis=1, inplace=True)
7 print("Train data size is : {}".format(df.shape))
8 print("Test data size is : {}".format(df_Test.shape))
9 print("Combined dataset size is : {}".format(all_data.shape))
```

Train data size is : (1157, 81)

Test data size is : (292, 80)

Combined dataset size is : (1449, 80)

```

In [41]: 1 # Find Missing Ratio of Dataset
2 all_data_na = (all_data.isnull().sum() / len(all_data)) * 100
3 all_data_na = all_data_na.drop(all_data_na[all_data_na == 0].index).sort_val
4 missing_data = pd.DataFrame({'Missing Ratio' :all_data_na})
5 missing_data

```


Out[41]:

	Missing Ratio
<b>PoolQC</b>	99.585921
<b>MiscFeature</b>	96.342305
<b>Alley</b>	93.788820
<b>Fence</b>	80.676329
<b>FireplaceQu</b>	47.412008
<b>LotFrontage</b>	17.805383
<b>GarageType</b>	5.590062
<b>GarageYrBlt</b>	5.590062
<b>GarageFinish</b>	5.590062
<b>GarageQual</b>	5.590062
<b>GarageCond</b>	5.590062
<b>BsmtExposure</b>	2.622498
<b>BsmtFinType2</b>	2.622498
<b>BsmtFinType1</b>	2.553485
<b>BsmtCond</b>	2.553485
<b>BsmtQual</b>	2.553485
<b>MasVnrArea</b>	0.552105
<b>MasVnrType</b>	0.552105
<b>Electrical</b>	0.069013



```
In [42]: 1 # Percent missing data by feature
2 f, ax = plt.subplots(figsize=(15, 12))
3 plt.xticks(rotation='90')
4 sns.barplot(x=all_data_na.index, y=all_data_na)
5 plt.xlabel('Features', fontsize=15)
6 plt.ylabel('Percent of missing values', fontsize=15)
7 plt.title('Percent missing data by feature', fontsize=15)
```

Out[42]: Text(0.5, 1.0, 'Percent missing data by feature')



```
1  Imputing Missing Values
2  PoolQC : data description says NA means "No Pool"
3  MiscFeature : data description says NA means "no misc feature"
4  Alley : data description says NA means "no alley access"
5  Fence : data description says NA means "no fence"
6  FireplaceQu : data description says NA means "no fireplace"
7  LotFrontage : Since the area of each street connected to the house property
   most likely have a similar area to other houses in its neighborhood , we
   can fill in missing values by the median LotFrontage of the neighborhood.
8  GarageType, GarageFinish, GarageQual and GarageCond : Replacing missing
   data with "None".
9  GarageYrBlt, GarageArea and GarageCars : Replacing missing data with 0.
10 BsmtFinSF1, BsmtFinSF2, BsmtUnfSF, TotalBsmtSF, BsmtFullBath and
   BsmtHalfBath: Replacing missing data with 0.
11 BsmtQual, BsmtCond, BsmtExposure, BsmtFinType1 and BsmtFinType2 : For all
   these categorical basement-related features, NaN means that there isn't a
   basement.
12 MasVnrArea and MasVnrType : NA most likely means no masonry veneer for
   these houses. We can fill 0 for the area and None for the type.
13 MSZoning (The general zoning classification) : 'RL' is by far the most
   common value. So we can fill in missing values with 'RL'.
14 Utilities : For this categorical feature all records are "AllPub", except
   for one "NoSeWa" and 2 NA . Since the house with 'NoSewa' is in the
   training set, this feature won't help in predictive modelling. We can then
   safely remove it.
15 Functional : data description says NA means typical.
16 Electrical : It has one NA value. Since this feature has mostly 'SBrkr', we
   can set that for the missing value.
17 KitchenQual: Only one NA value, and same as Electrical, we set 'TA' (which
   is the most frequent) for the missing value in KitchenQual.
18 Exterior1st and Exterior2nd : Both Exterior 1 & 2 have only one missing
   value. We will just substitute in the most common string
19 SaleType : Fill in again with most frequent which is "WD"
20 MSSubClass : Na most likely means No building class. We can replace missing
   values with None
```

```

In [43]: 1 all_data["PoolQC"] = all_data["PoolQC"].fillna("None")
2 all_data["MiscFeature"] = all_data["MiscFeature"].fillna("None")
3 all_data["Alley"] = all_data["Alley"].fillna("None")
4 all_data["Fence"] = all_data["Fence"].fillna("None")
5 all_data["FireplaceQu"] = all_data["FireplaceQu"].fillna("None")
6 all_data["LotFrontage"] = all_data.groupby("Neighborhood")["LotFrontage"].tr
7 for col in ('GarageType', 'GarageFinish', 'GarageQual', 'GarageCond'):
8     all_data[col] = all_data[col].fillna('None')
9 for col in ('GarageYrBlt', 'GarageArea', 'GarageCars'):
10     all_data[col] = all_data[col].fillna(0)
11 for col in ('BsmtFinSF1', 'BsmtFinSF2', 'BsmtUnfSF', 'TotalBsmtSF', 'BsmtFull
12 all_data[col] = all_data[col].fillna(0)
13 for col in ('BsmtQual', 'BsmtCond', 'BsmtExposure', 'BsmtFinType1', 'BsmtFin
14 all_data[col] = all_data[col].fillna('None')
15 all_data["MasVnrType"] = all_data["MasVnrType"].fillna("None")
16 all_data["MasVnrArea"] = all_data["MasVnrArea"].fillna(0)
17 all_data["MSZoning"] = all_data["MSZoning"].fillna(all_data["MSZoning"].mode
18 all_data = all_data.drop(['Utilities'], axis=1)
19 all_data["Functional"] = all_data["Functional"].fillna("Typ")
20 all_data["Electrical"] = all_data["Electrical"].fillna(all_data["Electrical"]
21 all_data["KitchenQual"] = all_data["KitchenQual"].fillna(all_data["KitchenQu
22 all_data["Exterior1st"] = all_data["Exterior1st"].fillna(all_data["Exterior1
23 all_data["Exterior2nd"] = all_data["Exterior2nd"].fillna(all_data["Exterior2
24 all_data["SaleType"] = all_data["SaleType"].fillna(all_data["SaleType"].mode
25 all_data["MSSubClass"] = all_data["MSSubClass"].fillna("None")

```

```

In [44]: 1 # Check if there are any missing values left
2 all_data_na = (all_data.isnull().sum() / len(all_data)) * 100
3 all_data_na = all_data_na.drop(all_data_na[all_data_na == 0].index).sort_val
4 missing_data = pd.DataFrame({'Missing Ratio' :all_data_na})
5 missing_data.head()

```

Out[44]:

Missing Ratio

## 5. Feature Transformation/Engineering

```
In [ ]: 1
2 Let's take a look at some features that may be misinterpreted to represent s
3
4 MSSubClass: Identifies the type of dwelling involved in the sale.
5
6 20 1-STORY 1946 & NEWER ALL STYLES
7 30 1-STORY 1945 & OLDER
8 40 1-STORY W/FINISHED ATTIC ALL AGES
9 45 1-1/2 STORY - UNFINISHED ALL AGES
10 50 1-1/2 STORY FINISHED ALL AGES
11 60 2-STORY 1946 & NEWER
12 70 2-STORY 1945 & OLDER
13 75 2-1/2 STORY ALL AGES
14 80 SPLIT OR MULTI-LEVEL
15 85 SPLIT FOYER
16 90 DUPLEX - ALL STYLES AND AGES
17 120 1-STORY PUD (Planned Unit Development) - 1946 & NEWER
18 150 1-1/2 STORY PUD - ALL AGES
19 160 2-STORY PUD - 1946 & NEWER
20 180 PUD - MULTILEVEL - INCL SPLIT LEV/FOYER
21 190 2 FAMILY CONVERSION - ALL STYLES AND AGES
```

```
In [45]: 1 all_data['MSSubClass'].describe()
```

```
Out[45]: count      1449.000000
mean         56.832298
std          42.277695
min          20.000000
25%          20.000000
50%          50.000000
75%          70.000000
max          190.000000
Name: MSSubClass, dtype: float64
```

So, the average is a 57 type. What does that mean? Is a 90 type 3 times better than a 30 type? This feature was interpreted as numerical when it is actually categorical. The types listed here are codes, not values. Thus, we need to feature transformation with this and many other features.

```
In [46]: 1 #MSSubClass =The building class
2 all_data['MSSubClass'] = all_data['MSSubClass'].apply(str)
3
4 #Changing OverallCond into a categorical variable
5 all_data['OverallCond'] = all_data['OverallCond'].astype(str)
6
7 #Year and month sold are transformed into categorical features.
8 all_data['YrSold'] = all_data['YrSold'].astype(str)
9 all_data['MoSold'] = all_data['MoSold'].astype(str)
```

```
1 Let's take a look at "Kitchen Quality".
```

```
In [47]: 1 all_data['KitchenQual'].unique()
```

```
Out[47]: array(['TA', 'Gd', 'Ex', 'Fa'], dtype=object)
```

Here, data\_description.txt comes to the rescue again!

Kitchen Quality:

Ex: Excellent Gd: Good TA: Typical/Average Fa: Fair Po: Poor Is a score of "Gd" better than "TA" but worse than "Ex"? I think so, let's encode these labels to give meaning to their specific orders.

```
In [48]: 1 from sklearn.preprocessing import LabelEncoder
2 cols = ('FireplaceQu', 'BsmtQual', 'BsmtCond', 'GarageQual', 'GarageCond',
3         'ExterQual', 'ExterCond', 'HeatingQC', 'PoolQC', 'KitchenQual', 'Bsmt
4         'BsmtFinType2', 'Functional', 'Fence', 'BsmtExposure', 'GarageFinish
5         'LotShape', 'PavedDrive', 'Street', 'Alley', 'CentralAir', 'MSSubCla
6         'YrSold', 'MoSold')
7 # Process columns and apply LabelEncoder to categorical features
8 for c in cols:
9     lbl = LabelEncoder()
10    lbl.fit(list(all_data[c].values))
11    all_data[c] = lbl.transform(list(all_data[c].values))
12
13 # Check shape
14 print('Shape all_data: {}'.format(all_data.shape))
```

Shape all\_data: (1449, 79)

Let's engineer one feature to combine square footage, this may be useful later on.

```
In [49]: 1 # Adding Total Square Feet feature
2 all_data['TotalSF'] = all_data['TotalBsmtSF'] + all_data['1stFlrSF'] + all_d
```

Fixing "skewed" features. Here, we fix all of the skewed data to be more normal so that our models will be more accurate when making predictions.


In [53]:


```

1  from scipy import stats
2  # We use the numpy function log1p which applies log(1+x) to all elements of
3  df["SalePrice"] = np.log1p(df["SalePrice"])
4
5
6  #Check the new distribution
7  sns.distplot(df['SalePrice'] , fit=norm);
8
9  # Get the fitted parameters used by the function
10 (mu, sigma) = norm.fit(df['SalePrice'])
11 print( '\n mu = {:.2f} and sigma = {:.2f}\n'.format(mu, sigma))
12 plt.legend(['Normal dist. ($\mu=${:.2f} and $\sigma=${:.2f} )'.format(mu,
13     loc='best')
14 plt.ylabel('Frequency')
15 plt.title('SalePrice distribution')
16
17 fig = plt.figure()
18 res = stats.probplot(df['SalePrice'], plot=plt)
19 plt.show()
20
21 y_train = df.SalePrice.values
22
23 print("Skewness: %f" % df['SalePrice'].skew())
24 print("Kurtosis: %f" % df['SalePrice'].kurt())

```

mu = 0.82 and sigma = 0.00





Skewness: -0.115359

Kurtosis: 1.199127

```
In [54]: 1 numeric_feats = all_data.dtypes[all_data.dtypes != "object"].index
          2
          3 # Check the skew of all numerical features
          4 skewed_feats = all_data[numeric_feats].apply(lambda x: skew(x.dropna())).sor
          5 skewness = pd.DataFrame({'Skewed Features' :skewed_feats})
          6 skewness.head()
```

Out[54]:

Skewed Features	
MiscVal	24.388024
PoolArea	15.882700
LotArea	12.595271
3SsnPorch	10.253854
LowQualFinSF	8.966866

```
In [55]: 1 skewness = skewness[abs(skewness) > 0.75]
          2 print("There are {} skewed numerical features to Box Cox transform".format(s
          3
          4 from scipy.special import boxcox1p
          5 skewed_features = skewness.index
          6 lam = 0.15
          7 for feat in skewed_features:
          8     all_data[feat] = boxcox1p(all_data[feat], lam)
          9     all_data[feat] += 1
```

There are 60 skewed numerical features to Box Cox transform

```
In [56]: 1 all_data = pd.get_dummies(all_data)
          2 print(all_data.shape)
```

(1449, 221)

```
In [57]: 1 train = all_data[:ntrain]
2 test = all_data[ntrain:]
```

## 6. Modeling and Predictions

```
In [59]: 1 from sklearn.linear_model import ElasticNet, Lasso, BayesianRidge, LassoLar
2 from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegress
3 from sklearn.kernel_ridge import KernelRidge
4 from sklearn.pipeline import make_pipeline
5 from sklearn.preprocessing import RobustScaler
6 from sklearn.base import BaseEstimator, TransformerMixin, RegressorMixin, cl
7 from sklearn.model_selection import KFold, cross_val_score, train_test_split
8 from sklearn.metrics import mean_squared_error
9
```

```
In [60]: 1 # Cross-validation with k-folds
2 n_folds = 5
3
4 def rmsle_cv(model):
5     kf = KFold(n_folds, shuffle=True, random_state=42).get_n_splits(train.va
6     rmse = np.sqrt(-cross_val_score(model, train.values, y_train, scoring="ne
7     return(rmse)
```

For our models, we are going to use lasso, elastic net, kernel ridge, gradient boosting

```
In [61]: 1 lasso = make_pipeline(RobustScaler(), Lasso(alpha =0.0005, random_state=1))
2 ENet = make_pipeline(RobustScaler(), ElasticNet(alpha=0.0005, l1_ratio=.9, r
3 KRR = KernelRidge(alpha=0.6, kernel='polynomial', degree=2, coef0=2.5)
4 GBoost = GradientBoostingRegressor(n_estimators=3000, learning_rate=0.05,
5                                     max_depth=4, max_features='sqrt',
6                                     min_samples_leaf=15, min_samples_split=10
7                                     loss='huber', random_state =5)
```



```
In [62]: 1 score = rmsle_cv(lasso)
2 print("\nLasso score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))
3 score = rmsle_cv(ENet)
4 print("ElasticNet score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))
5 score = rmsle_cv(KRR)
6 print("Kernel Ridge score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))
7 score = rmsle_cv(GBoost)
8 print("Gradient Boosting score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))
9
```

Lasso score: 0.0017 (0.0003)

ElasticNet score: 0.0017 (0.0003)

Kernel Ridge score: 0.0019 (0.0003)

Gradient Boosting score: 0.0012 (0.0002)

Here, we stack the models to average their scores.

```
In [63]: 1 class AveragingModels(BaseEstimator, RegressorMixin, TransformerMixin):
2     def __init__(self, models):
3         self.models = models
4
5         # we define clones of the original models to fit the data in
6     def fit(self, X, y):
7         self.models_ = [clone(x) for x in self.models]
8
9         # Train cloned base models
10        for model in self.models_:
11            model.fit(X, y)
12
13        return self
14
15        #Now we do the predictions for cloned models and average them
16    def predict(self, X):
17        predictions = np.column_stack([
18            model.predict(X) for model in self.models_
19        ])
20        return np.mean(predictions, axis=1)
```

Here we average ENet, GBoost, KRR, and lasso

```
In [64]: 1 averaged_models = AveragingModels(models = (ENet, GBoost, KRR, lasso))
2
3 score = rmsle_cv(averaged_models)
4 print("Averaged base models score: {:.4f} ({:.4f})\n".format(score.mean(), score.std()))
```

Averaged base models score: 0.0013 (0.0002)

```

In [65]: 1 class StackingAveragedModels(BaseEstimator, RegressorMixin, TransformerMixin)
2         def __init__(self, base_models, meta_model, n_folds=5):
3             self.base_models = base_models
4             self.meta_model = meta_model
5             self.n_folds = n_folds
6
7         # We again fit the data on clones of the original models
8         def fit(self, X, y):
9             self.base_models_ = [list() for x in self.base_models]
10            self.meta_model_ = clone(self.meta_model)
11            kfold = KFold(n_splits=self.n_folds, shuffle=True)
12
13            # Train cloned base models then create out-of-fold predictions
14            # that are needed to train the cloned meta-model
15            out_of_fold_predictions = np.zeros((X.shape[0], len(self.base_models)
16            for i, clf in enumerate(self.base_models):
17                for train_index, holdout_index in kfold.split(X, y):
18                    instance = clone(clf)
19                    self.base_models_[i].append(instance)
20                    instance.fit(X[train_index], y[train_index])
21                    y_pred = instance.predict(X[holdout_index])
22                    out_of_fold_predictions[holdout_index, i] = y_pred
23
24            # Now train the cloned meta-model using the out-of-fold predictions
25            self.meta_model_.fit(out_of_fold_predictions, y)
26            return self
27
28        def predict(self, X):
29            meta_features = np.column_stack([
30                np.column_stack([model.predict(X) for model in base_models]).mean(
31                    for base_models in self.base_models_ ])
32            return self.meta_model_.predict(meta_features)

```

Since our lasso model performed the best, we'll use it as a meta-model.

```

In [66]: 1 stacked_averaged_models = StackingAveragedModels(base_models = (ENet, GBoost
2                                     meta_model = lasso)
3
4         score = rmsle_cv(stacked_averaged_models)
5         print("Stacking Averaged models score: {:.4f} ({:.4f})".format(score.mean(),

```

Stacking Averaged models score: 0.0013 (0.0002)

```

In [67]: 1 def rmsle(y, y_pred):
2         return np.sqrt(mean_squared_error(y, y_pred))

```

```

1 Stacked models

```

```
In [68]: 1 stacked_averaged_models.fit(train.values, y_train)
          2 stacked_train_pred = stacked_averaged_models.predict(train.values)
          3 stacked_pred = np.expm1(stacked_averaged_models.predict(test.values))
          4 print(rmsle(y_train, stacked_train_pred))
```

0.0009936753449509559

```
1 Conclusion : Since our Lasso model performs well we accept this model
```

submission

```
In [71]: 1 sub = pd.DataFrame()
          2 sub['Id'] = df_Test.Id
          3 sub['SalePrice'] = stacked_pred
          4 sub.to_csv('submission.csv',index=False)
```

```
In [ ]: 1
```